



# Prototype networks for few-shot learning

Fabian Greavu, Machine Learning exam @ unifi



# Summary

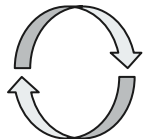
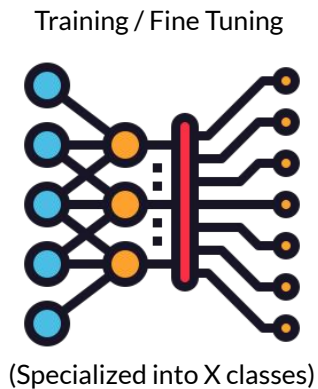
- ❖ Introduction
  - Challenge
  - Few-shot learning
- ❖ Project workflow
  - Datasets
  - Prototypical Networks
  - Centroids
  - NC, NS, NQ
  - Loss, Optimizer, scheduler
  - Training skeleton
  - Code
- ❖ Experiments
  - Omniglot
  - Mini Imagenet
  - Flowers102
  - Comparison
- ❖ Conclusion





# Intro

# Challenge

'classic' Classification Task




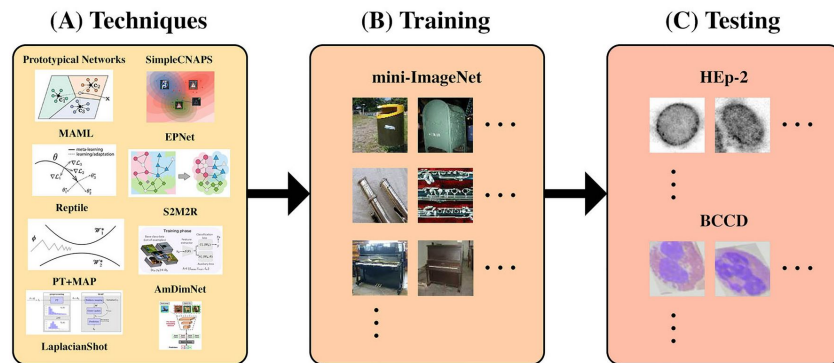
Our Challenge Classification Task



...

# Few-shot learning

- A **meta-learning** technique to 'learn to learn'
- Uses **small amount** of occurrences
- Can perform with **high accuracies**



# Meta-learning definition

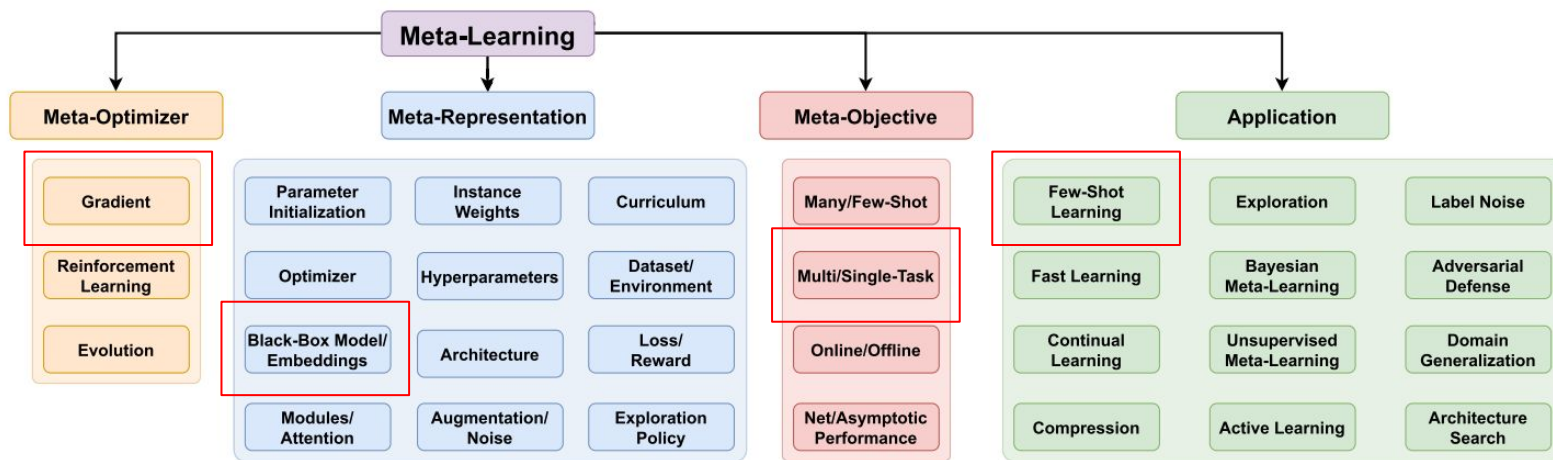


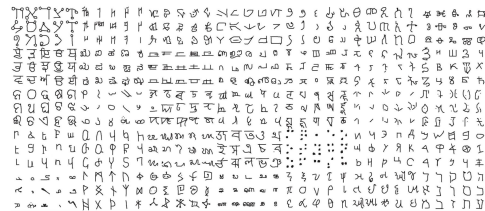
Fig. 1. Overview of the meta-learning landscape including algorithm design (meta-optimizer, meta-representation, meta-objective), and applications.



# Project workflow

# Datasets

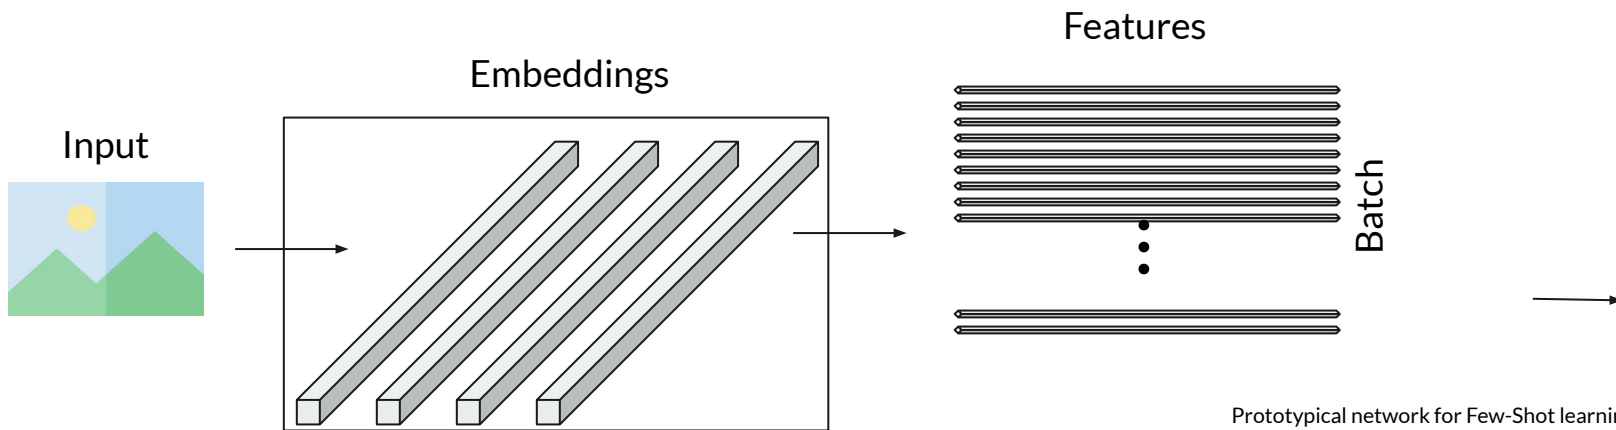
- Omniglot: 1623 handwritten characters
  - 80 images per class
  - classes: (1032 train, 172 val, 464 test)
- Mini Imagenet: 100 objects
  - 600 images per class
  - classes: (64 train, 16 val, 20 test)
- Flowers102: 102 flowers
  - 40-120 images per class
  - classes: (64 train, 16 val, 22 test)





# Prototypical Networks

- Extracts features with a neural network
  - 4 CNN blocks: Conv out=64, ks=3 / BatchNorm2D / ReLu / MaxPool2D
- Learning: embeddings learning
  - Trained 100 episodes/iterations per epochs (200 ep) to learn embeddings



# Prototypical Networks - centroids

- Use a small part of output as support, other as query
- Centroids are mean(support)
- Calculate distances between query and centroids
- Calculate loss as mean of log\_softmax of negative distances


Prototypical networks compute an  $M$ -dimensional representation  $\mathbf{c}_k \in \mathbb{R}^M$ , or *prototype*, of each class through an embedding function  $f_\phi : \mathbb{R}^D \rightarrow \mathbb{R}^M$  with learnable parameters  $\phi$ . Each prototype is the mean vector of the embedded support points belonging to its class:

$$\mathbf{c}_k = \frac{1}{|S_k|} \sum_{(\mathbf{x}_i, y_i) \in S_k} f_\phi(\mathbf{x}_i) \quad (1)$$

Given a distance function  $d : \mathbb{R}^M \times \mathbb{R}^M \rightarrow [0, +\infty)$ , prototypical networks produce a distribution over classes for a query point  $\mathbf{x}$  based on a softmax over distances to the prototypes in the embedding space:

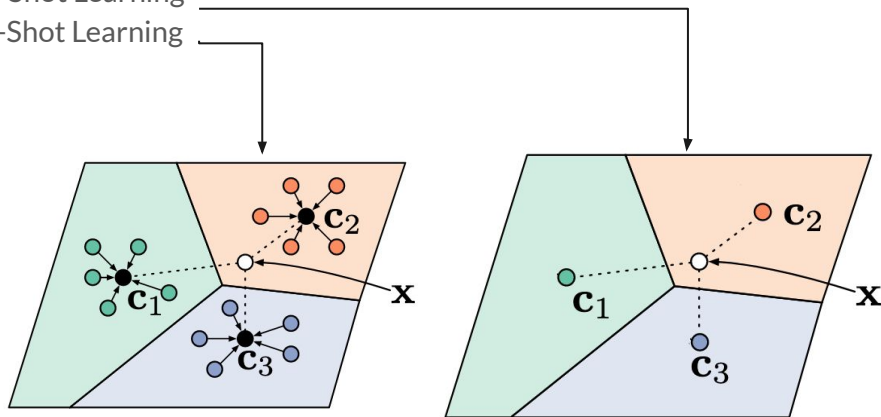
$$p_\phi(y = k | \mathbf{x}) = \frac{\exp(-d(f_\phi(\mathbf{x}), \mathbf{c}_k))}{\sum_{k'} \exp(-d(f_\phi(\mathbf{x}), \mathbf{c}_{k'}))} \quad (2)$$

Learning proceeds by minimizing the negative log-probability  $J(\phi) = -\log p_\phi(y = k | \mathbf{x})$  of the true class  $k$  via SGD. Training episodes are formed by randomly selecting a subset of classes from the training set, then choosing a subset of examples within each class to act as the support set and a subset of the remainder to serve as query points. Pseudocode to compute the loss  $J(\phi)$  for a training episode is provided in Algorithm 1.



# Prototypical Networks - NC, NS, NQ

- NC: how many classes to use per each iteration on batch (or ‘ways’)
- NS: how many examples to use as support for centroids calculus (or ‘shots’)
- NQ: how many examples to use as queries for centroids calculus (‘query’)
  - 1 shots -> One-Shot Learning
  - 5 shots -> Few-Shot Learning





## Loss, optimizer, scheduler

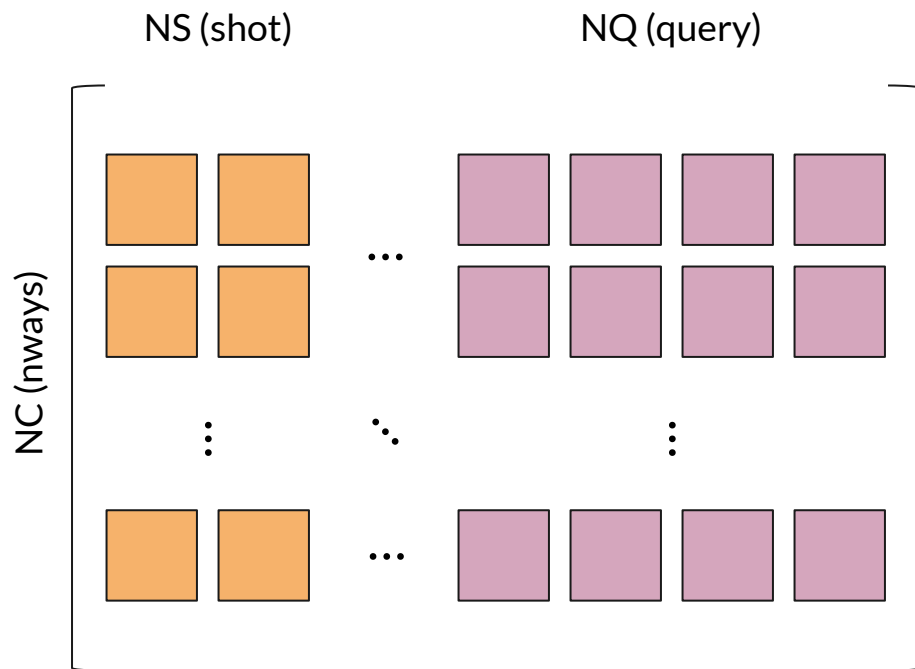
- Loss: `mean( log_softmax( -distances( query, centroids ), targets )`
- Optimizer: Adam with `lr=0.001` (faster and similar results to SGD)
- Scheduler: StepLr with `step_size=20, gamma = 0.5`

# Training skeleton

```
model = PrototypicalNet()
loss = loss_function(x, y, NC, NS, NQ)
optim = StepLr()
```

```
for epoch in epochs:
    for it in iterations:
        x, y = GetSample(NC, NS, NQ)
        out = model(x)
        ... loss
        backward()
    for it in iterations:
        .... eval ....
```

One batch





# Code

All code is available at [github](#)

- Full hyperparams control
- 3 available datasets
- results and plots
- simple train.py and test.py scripts

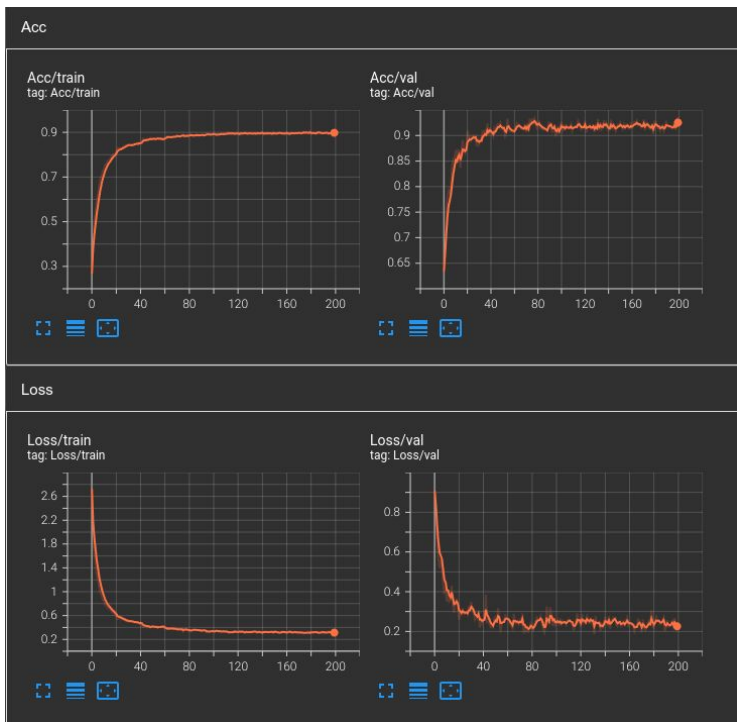
```
python train.py --dataset mini_imagenet \  
    --epochs 200 \  
    --gpu \  
    --train-num-class 30 \  
    --test-num-class 5 \  
    --number-support 5 \  
    --train-num-query 15 \  
    --episodes-per-epoch 100 \  
    --adam-lr 0.001 \  
    --opt-step-size 20 \  
    --opt-gamma 0.5 \  
    --distance-function "euclidean" \  
    --save-each 5
```



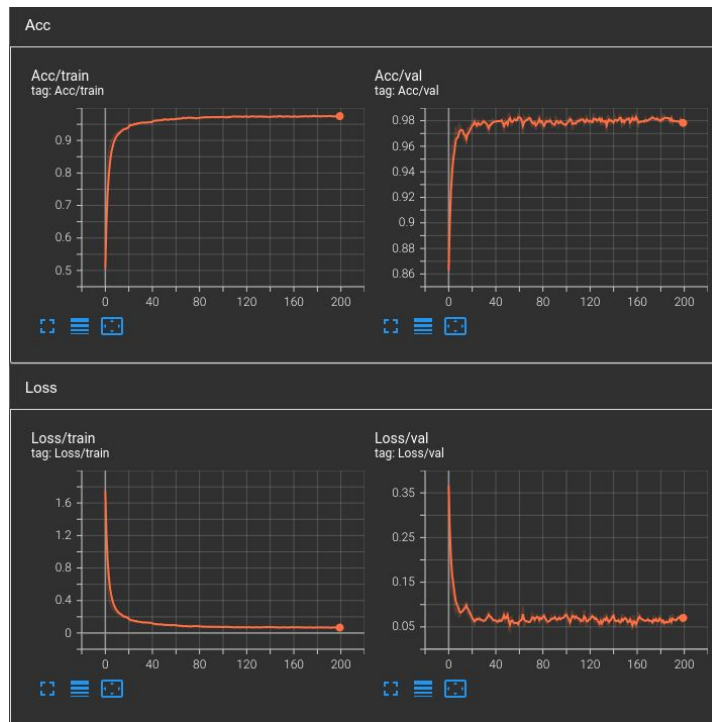
# Experiments

# Omniglot

1-shot



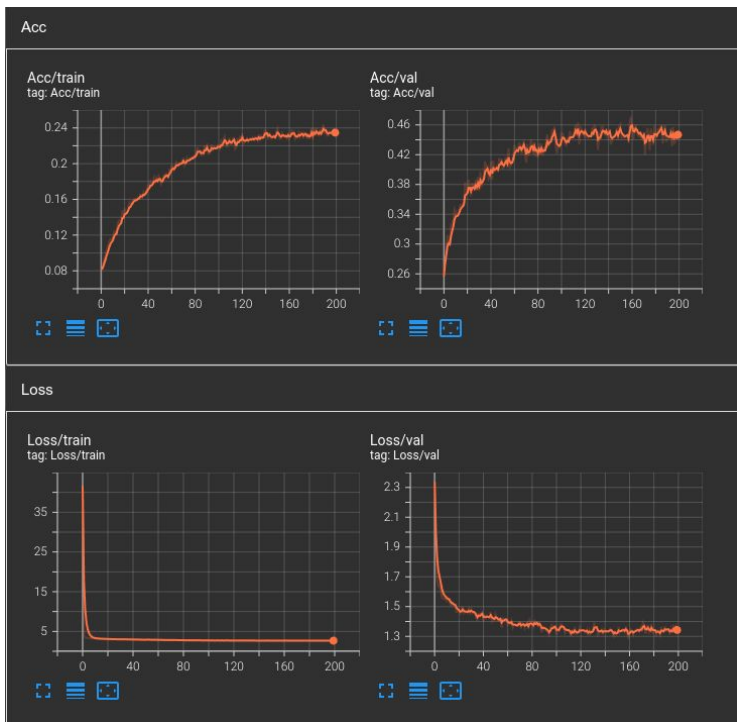
5-shot



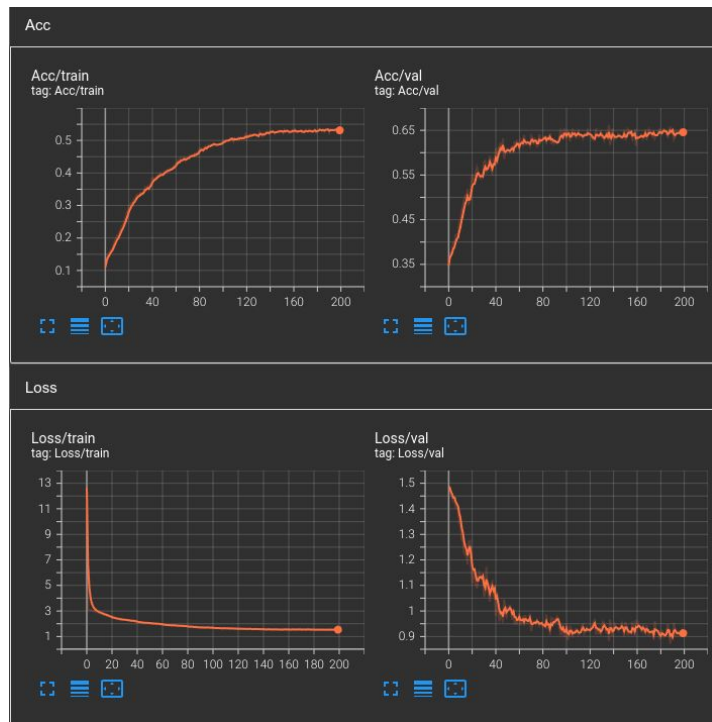


# Mini Imagenet

1-shot

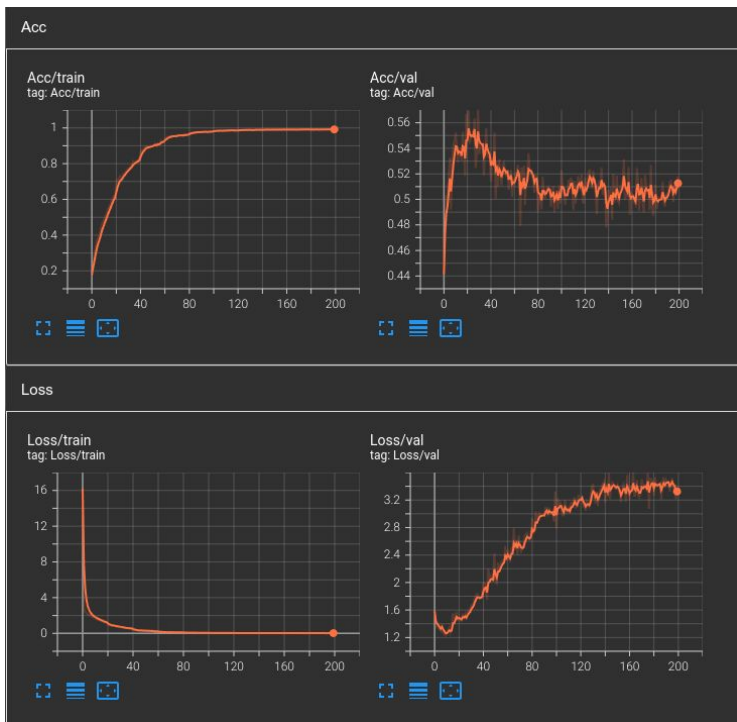


5-shot

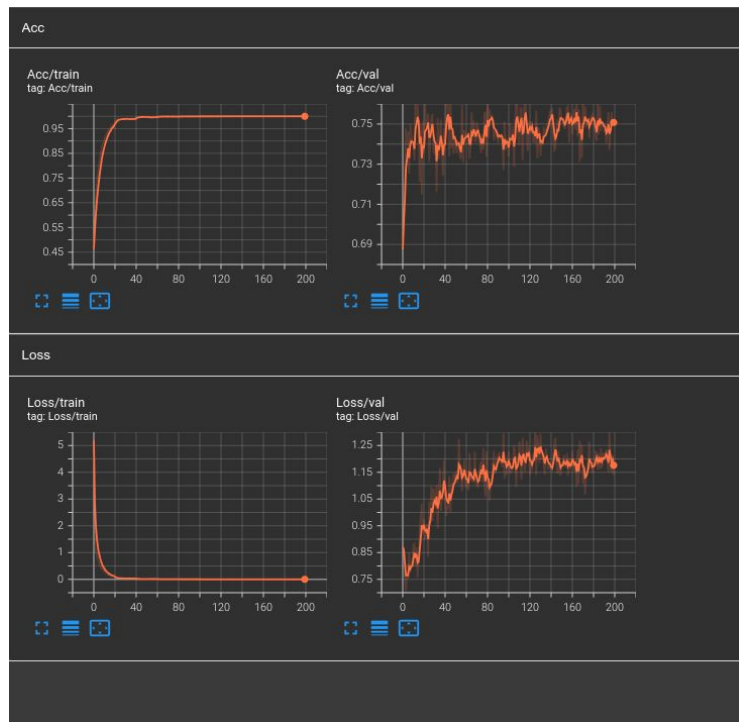


# Flowers 102

1-shot



5-shot





## Comparison - distance metric

<b>Dataset</b>	<b>Cosine (acc)</b>	<b>Euclidean (acc)</b>
mini_imagenet	22.36	<b>63.62</b>
omniglot	23.48	<b>97.77</b>
flowers102	82.89	<b>84.48</b>



## Comparison - one vs few shot vs paper

Dataset	Paper res 5-way 5-shot (Acc)	Our res 5-way 5-shot (Acc)	Paper res 5-way 1-shot (Acc)	Our res 5-way 1-shot (Acc)
mini_imagenet	68.20	63.62	49.42	46.13
omniglot	98.80	97.77	98.8	91.93
flowers102	/	84.48	/	56.08



# Conclusions



## Conclusion

- Euclidean distance performs better than cosine similarity
- Paper results were correctly replicated

## Future studies

- Add custom dataset option for training
- Implement proper `torch.nn.Dataset` and `torch.nn.Sampler` + `torch.nn.DataLoader`
- Try different fields than CV